

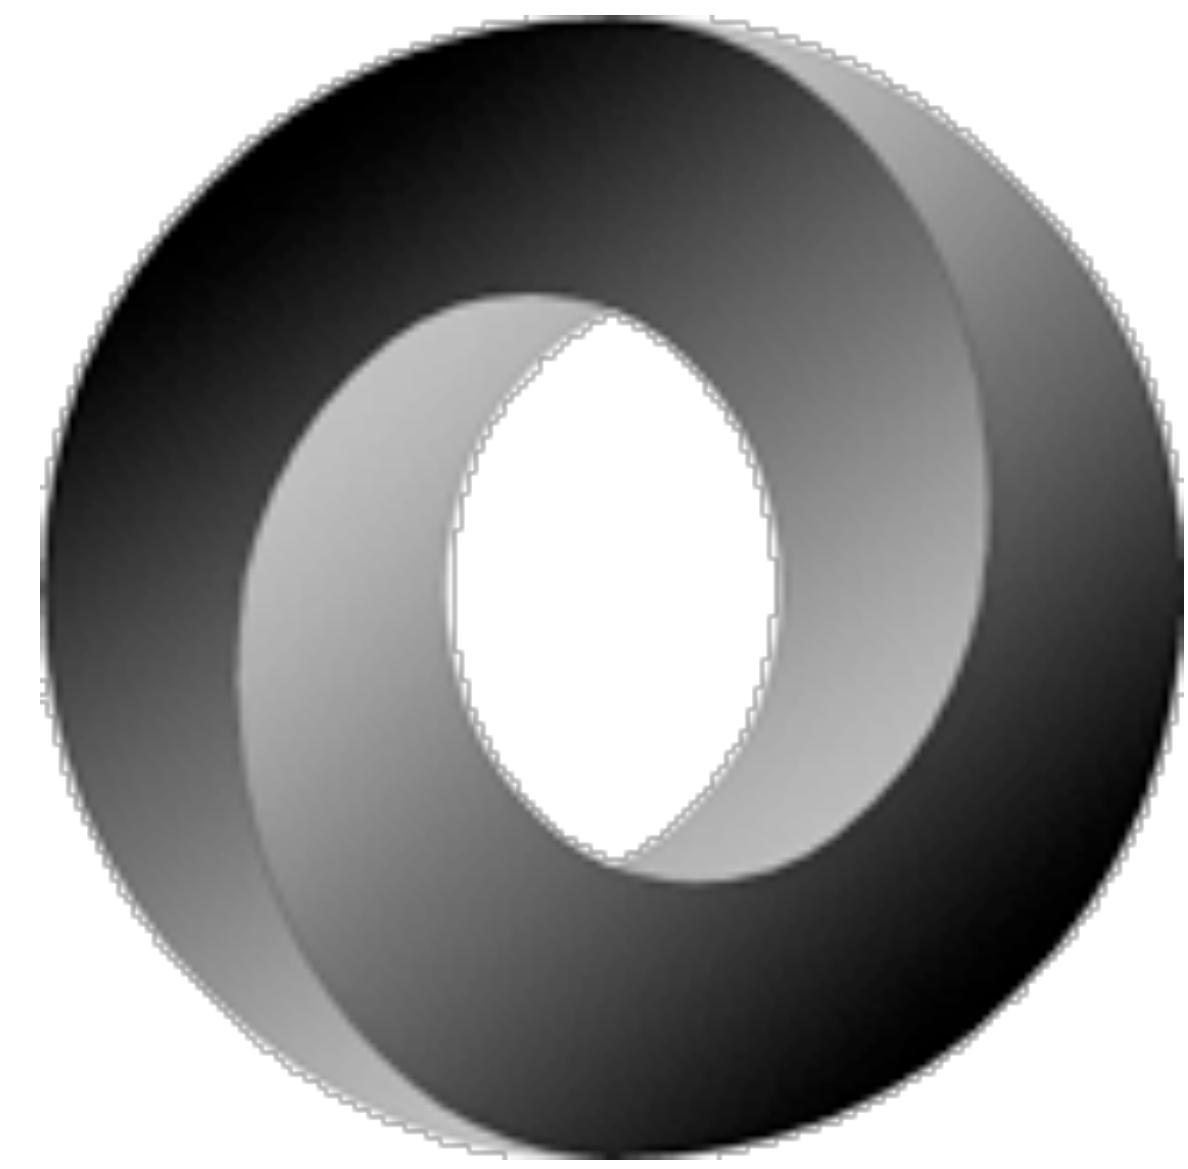
Future JSON support

looking ahead...

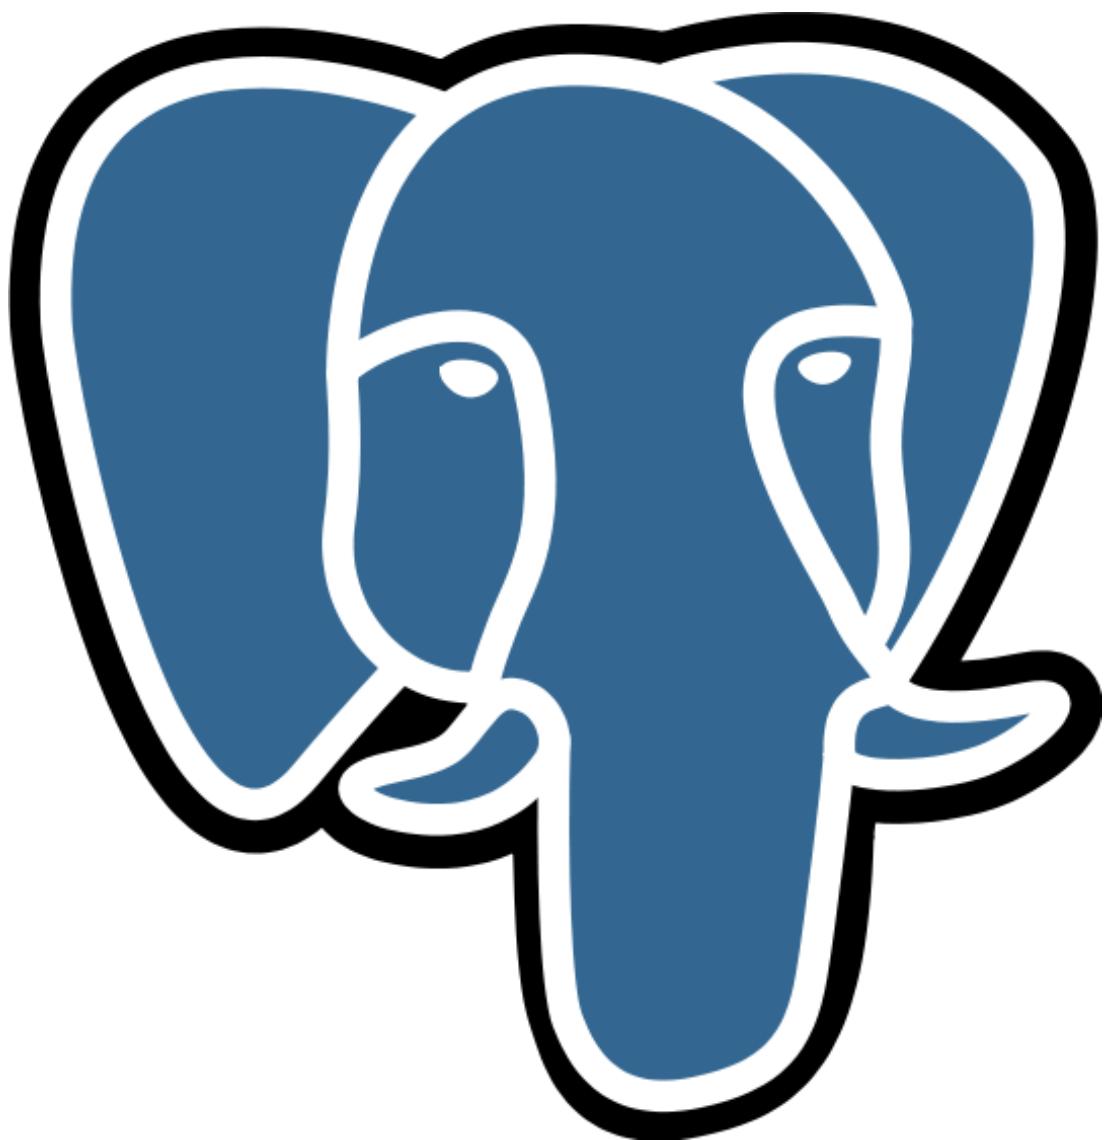


Daniel Gustafsson

Citus Team Meeting * 31/7 2023



+



=



JSON in PostgreSQL

JSON type in 9.2

JSONB type in 9.4

New functions and optimizations in every release like FTS support in 10.0 and SQL/JSON PATH language in 12.0

Widely appreciated. Commonly reported as faster than MongoDB on JSON for data load and queries using an index

SQL/JSON support coming in 16.0



17



BIKESHED

ADVISORY

WARNING

```
[  
  {  
    "item": "kiwi",  
    "inv": 10  
  },  
  {  
    "item": "pear",  
    "inv": 20  
  }  
]
```

```
CREATE TABLE inventory
(
    products jsonb
);
```

```
INSERT INTO inventory VALUES
(
    '[
        {"item":"kiwi", "inv":10},
        {"item":"pear", "inv":20}
    ]'
);
```

```
SELECT
    jsonb_array_elements(products)->>'item'
        AS item,
    jsonb_array_elements(products)->>'inv'
        AS inventory
FROM
    inventory;
```

item	inventory
kiwi	10
pear	20

(2 rows)

```
{ "produce": [
    { "type": "fruits",
      "items": [
          { "item": "kiwi", "inv": 10 },
          { "item": "pear", "inv": 20 }
      ]
    },
    { "type": "vegetables",
      "items": [
          { "item": "lettuce", "inv": 100 },
          { "item": "celery", "inv": 20 }
      ]
    }
  ]
}
```





stack overflow



```
SELECT
    e.v->>'type' AS type,
    jsonb_array_elements(e.v->'items')->>'item'
        AS item,
    jsonb_array_elements(e.v->'items')->>'inv'
        AS inventory
FROM
    inventory
    CROSS JOIN lateral jsonb_each(products)
        AS p(k,v)
    CROSS JOIN lateral jsonb_array_elements(p.v)
        AS e(v);
```

type	item	inventory
fruits	kiwi	10
fruits	pear	20
vegetables	lettuce	100
vegetables	celery	20
(4 rows)		

QUERY PLAN

Result

-> ProjectSet

-> Nested Loop

-> Nested Loop

-> Seq Scan on inventory

-> Function Scan on jsonb_each p

-> Memoize

Cache Key: p.v

Cache Mode: binary

-> Function Scan on jsonb_array_elements e

(10 rows)

```
SELECT
    e.v->>'type' AS type,
    jsonb_array_elements(e.v->'items')->>'item'
        AS item,
    jsonb_array_elements(e.v->'items')->>'inv'
        AS inventory
FROM
    inventory
    CROSS JOIN lateral jsonb_each(products)
        AS p(k,v)
    CROSS JOIN lateral jsonb_array_elements(p.v)
        AS e(v);
```

```
SELECT
    e.v->>'type' AS type,
    jsonb_array_elements(e.v->'items')->>'item'
        AS item,
    jsonb_array_elements(e.v->'items')->>'inv'
        AS inventory
FROM
    inventory
    CROSS JOIN lateral jsonb_each(products)
        AS p(k,v)
    CROSS JOIN lateral jsonb_array_elements(p.v)
        AS e(v);
```

7.11 <JSON table>

Function

Query a JSON text and present it as a relational table.

Format

```
<JSON table> ::=  
    JSON_TABLE <left paren>  
        <JSON API common syntax>  
        <JSON table columns clause>  
        [ <JSON table plan clause> ]  
        [ <JSON table error behavior> ON ERROR ]  
    <right paren>  
  
<JSON table columns clause> ::=  
    COLUMNS <left paren>  
        <JSON table column definition> [ { <comma> <JSON table column definition> }... ]  
    <right paren>  
  
<JSON table column definition> ::=  
    <JSON table ordinality column definition>  
    | <JSON table no ordinality column definition>
```

JSON_TABLE

Standardized SQL

Supported by IBM DB2, Oracle and MySQL

Query JSONB structures and present as relational data

Resulting relation can be treated as ordinary relation content

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
));

```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
))
);

```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        {
          "item": "kiwi", "inv": 10
        },
        {
          "item": "pear", "inv": 20
        }
      ]
    },
    {
      "type": "vegetables",
      "items": [
        {
          "item": "lettuce", "inv": 100
        },
        {
          "item": "celery", "inv": 20
        }
      ]
    }
  ]
}
```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
)
)
);
```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        {
          "item": "kiwi", "inv": 10
        },
        {
          "item": "pear", "inv": 20
        }
      ]
    },
    {
      "type": "vegetables",
      "items": [
        {
          "item": "lettuce", "inv": 100
        },
        {
          "item": "celery", "inv": 20
        }
      ]
    }
  ]
}
```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
))
);

```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        {
          "item": "kiwi", "inv": 10
        },
        {
          "item": "pear", "inv": 20
        }
      ]
    },
    {
      "type": "vegetables",
      "items": [
        {
          "item": "lettuce", "inv": 100
        },
        {
          "item": "celery", "inv": 20
        }
      ]
    }
  ]
}
```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
)
) );
```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        {
          "item": "kiwi", "inv": 10
        },
        {
          "item": "pear", "inv": 20
        }
      ]
    },
    {
      "type": "vegetables",
      "items": [
        {
          "item": "lettuce", "inv": 100
        },
        {
          "item": "celery", "inv": 20
        }
      ]
    }
  ]
}
```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
)
);
```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        { "item": "kiwi", "inv": 10 },
        { "item": "pear", "inv": 20 }
      ]
    },
    {
      "type": "vegetables",
      "items": [
        { "item": "lettuce", "inv": 100 },
        { "item": "celery", "inv": 20 }
      ]
    }
  ]
}
```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
)
) );
```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        {"item": "kiwi", "inv": 10},
        {"item": "pear", "inv": 20}
      ]
    },
    {
      "type": "vegetables",
      "items": [
        {"item": "lettuce", "inv": 100},
        {"item": "celery", "inv": 20}
      ]
    }
  ]
}
```

```
SELECT
produce,
item,
inv
FROM
inventory,
json_table(products, '$.produce[*]' COLUMNS (
produce TEXT path '$.type',
NESTED PATH '$.items[*]' COLUMNS (
item TEXT path '$.item',
inv integer path '$.inv'
)
)
) );
```

```
{
  "produce": [
    {
      "type": "fruits",
      "items": [
        {
          "item": "kiwi", "inv": 10
        },
        {
          "item": "pear", "inv": 20
        }
      ]
    },
    {
      "type": "vegetables",
      "items": [
        {
          "item": "lettuce", "inv": 100
        },
        {
          "item": "celery", "inv": 20
        }
      ]
    }
  ]
}
```

type	item	inventory
fruits	kiwi	10
fruits	pear	20
vegetables	lettuce	100
vegetables	celery	20
(4 rows)		

QUERY PLAN

Nested Loop

-> Seq Scan on inventory
-> Table Function Scan on "json_table"

(3 rows)

```
SELECT
    produce,
    sum(inv) AS total_inventory
FROM
    inventory,
    json_table(products, '$.produce[*]' COLUMNS (
        produce TEXT path '$.type',
        NESTED PATH '$.items[*]' COLUMNS (
            item TEXT path '$.item',
            inv integer path '$.inv'
        )
    ))
GROUP BY produce;
```

produce	total_inventory
vegetables	120
fruits	30
(2 rows)	

```
SELECT  
    id,  
    produce,  
    item,  
    inv  
FROM  
    inventory,  
    json_table(products, '$.produce[*]' COLUMNS (  
        id FOR ORDINALITY,  
        produce TEXT path '$.type',  
        NESTED PATH '$.items[*]' COLUMNS (  
            item TEXT path '$.item',  
            inv integer path '$.inv'  
        )  
    ));
```

id	produce	item	inv
1	fruits	kiwi	10
1	fruits	pear	20
2	vegetables	lettuce	100
2	vegetables	celery	20

(4 rows)

```
SELECT
*
FROM
json_table(
  '[{"item":"kiwi","inv":10},
   {"item":"pear","inv":20}]' ::JSONB,
  '$[*]' COLUMNS(
    item TEXT path '$.item',
    inventory INTEGER path '$.inv'
  )
);
```

TL;DR - JSON_TABLE

Returns a relation like any other

Advanced SQL analytics

Not limited to a 1:1 mapping

Transformative declaration

name type EXISTS ...

name type FORMAT ...

JSON_QUERY(*context*, *path* ...)

Returns the result of applying path expression to the context expression using a set of optional values.

Always returns a JSON string

Can return multiple objects

Expressive error semantics for controlling missing or incorrect data, ERROR, EMPTY, NULL, DEFAULT

```
SELECT
    json_query(products, '$.produce[*].items[1].item'
    WITH WRAPPER)
FROM
    inventory;
```

json_query

```
-----
["pear", "celery"]
(1 row)
```

JSON_VALUE(*context*, *path* ...)

Returns the result of applying path expression to the context expression using a set of optional values.

Can only return a single SQL/JSON scalar item

Expressive error semantics for controlling missing or incorrect data, ERROR, EMPTY, NULL, DEFAULT

```
SELECT
    json_value(jsonb '"03:04 2015-02-01"',
    '$.datetime("HH24:MI YYYY-MM-DD")'
RETURNING date);
```

json_value

2015-02-01
(1 row)

JSON_EXISTS(*context*, *path* ...)

Returns the boolean True if applying the path expression to the context expression with optional values yields any results.

```
SELECT
    json_exists(products, '$.produce[*].items')
FROM
    inventory;
```

json_exists

t
(1 row)



SQL/JSON Constructors

json_array()

json_arrayagg()

json_object()

json_objectagg()

```
SELECT  
    json_array(SELECT relname FROM pg_class  
              ORDER BY oid ASC limit 3);
```

json_array

```
-----  
[ "pg_foreign_data_wrapper_oid_index",  
  "pg_foreign_server_oid_index",  
  "pg_user_mapping_oid_index"]  
(1 row)
```

```
SELECT  
    json_object('fruit' value 'pear',  
                'vegetable' : 'celery');
```

json_object

```
{"fruit" : "pear", "vegetable" : "celery"}  
(1 row)
```

```
SELECT  
    json_object('fruit' value 'pear',  
                'vegetable' : 'celery');
```

json_object

```
{"fruit" : "pear", "vegetable" : "celery"}  
(1 row)
```

SQL/JSON object checks

expression IS [NOT] JSON ...

Tests if expression can be parsed as JSON, possibly of a specified type

Support for values, arrays, objects, scalars and unique keys

```
SELECT
  ' ' IS JSON AS underscore,
  '{}' IS JSON AS brackets;
```

underscore	brackets
f	t

(1 row)

SQL/JSON Summary

Constructors and vtests coming in v16

JSON_TABLE and expression based extraction functionality
hopefully coming in v17

Some functionality has already been committed



Thats all folks!